

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
```

```
In [2]: df = pd.read_csv("Bangalore Housing Prices.csv")
```

```
In [4]: df.head()
```

```
Out[4]:
```

	location	size	total_sqft	bath	price
0	Electronic City Phase II	2 BHK	1056	2.0	39.07
1	Chikka Tirupathi	4 Bedroom	2600	5.0	120.00
2	Uttarahalli	3 BHK	1440	2.0	62.00
3	Lingadheeranahalli	3 BHK	1521	3.0	95.00
4	Kothanur	2 BHK	1200	2.0	51.00

```
In [5]: df.shape
```

```
Out[5]: (13320, 5)
```

```
In [6]: df.describe
```

```
Out[6]: <bound method NDFrame.describe of
ice
```

	location	size	total_sqft	bath	pr
0	Electronic City Phase II	2 BHK	1056	2.0	39.07
1	Chikka Tirupathi	4 Bedroom	2600	5.0	120.00
2	Uttarahalli	3 BHK	1440	2.0	62.00
3	Lingadheeranahalli	3 BHK	1521	3.0	95.00
4	Kothanur	2 BHK	1200	2.0	51.00
...	...	...	...	...	...
13315	Whitefield	5 Bedroom	3453	4.0	231.00
13316	Richards Town	4 BHK	3600	5.0	400.00
13317	Raja Rajeshwari Nagar	2 BHK	1141	2.0	60.00
13318	Padmanabhanagar	4 BHK	4689	4.0	488.00
13319	Doddathoguru	1 BHK	550	1.0	17.00

```
[13320 rows x 5 columns]>
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13320 entries, 0 to 13319
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   location    13319 non-null  object
1   size        13304 non-null  object
2   total_sqft  13320 non-null  object
3   bath        13247 non-null  float64
```

```
4    price    13320 non-null    float64
dtypes: float64(2), object(3)
memory usage: 520.4+ KB
```

```
In [8]: df.isnull().sum()
```

```
Out[8]: location      1
size      16
total_sqft    0
bath      73
price      0
dtype: int64
```

```
In [9]: df.dropna(inplace = True)
```

```
In [12]: df.isnull().sum()
```

```
Out[12]: location      0
size      0
total_sqft    0
bath      0
price      0
dtype: int64
```

```
In [13]: df['size'] = [int(value.split(' ')[0]) for value in df['size']]
```

```
In [14]: df.head()
```

```
Out[14]:
```

	location	size	total_sqft	bath	price
0	Electronic City Phase II	2	1056	2.0	39.07
1	Chikka Tirupathi	4	2600	5.0	120.00
2	Uttarahalli	3	1440	2.0	62.00
3	Lingadheeranahalli	3	1521	3.0	95.00
4	Kothanur	2	1200	2.0	51.00

```
In [22]: df['total_sqft'].describe
```

```
Out[22]: <bound method NDFrame.describe of 0      1056
1      2600
2      1440
3      1521
4      1200
...
13315   3453
13316   3600
13317   1141
13318   4689
13319    550
Name: total_sqft, Length: 13246, dtype: object>
```

```
In [23]: def convert_sqft(value):
          try:
              if '-' in value:
                  start, end = map(float, value.split('-'))
```

```

        return (start + end) / 2
    else:
        return float(value)
except ValueError:
    return float('nan')

```

```

In [24]: df['total_sqft'] = [convert_sqft(value) for value in df['total_sqft']]

```

```

In [25]: df['total_sqft']

```

```

Out[25]: 0      1056.0
         1      2600.0
         2      1440.0
         3      1521.0
         4      1200.0
         ...
        13315    3453.0
        13316    3600.0
        13317     1141.0
        13318    4689.0
        13319     550.0
        Name: total_sqft, Length: 13246, dtype: float64

```

```

In [26]: df['total_sqft'].isnull().sum()

```

```

Out[26]: 46

```

```

In [27]: df.isnull().sum()

```

```

Out[27]: location      0
         size          0
         total_sqft    46
         bath          0
         price          0
         dtype: int64

```

```

In [28]: df.dropna(inplace = True)

```

```

In [29]: df.isnull().sum()

```

```

Out[29]: location      0
         size          0
         total_sqft    0
         bath          0
         price          0
         dtype: int64

```

```

In [30]: df['Price_per_sqft'] = df['price']/df['total_sqft']

```

```



In [31]: selected_columns = ['Price_per_sqft', 'size']
         outliers = df[selected_columns]

```

```

In [33]: for i in outliers:
         sns.boxplot(x=df[i])
         plt.show()


```




```
In [34]: def remove_outliers(column):
          Q1 = column.quantile(0.25)
          Q3 = column.quantile(0.75)
          IQR = Q3 - Q1
          lower_bound = Q1 - 1.5 * IQR
          upper_bound = Q3 + 1.5 * IQR
          return column[(column >= lower_bound) & (column <= upper_bound)]
```

```
In [35]: df['Price_per_sqft'] = remove_outliers(df['Price_per_sqft'])
```

```
In [36]: sns.boxplot(x=df['Price_per_sqft'])
          plt.show()
```



```
In [37]: sns.boxplot(x=df['Price_per_sqft'])
plt.show()
```



```
In [39]: df.isnull().sum()
```

```
Out[39]: location          0
size          0
total_sqft      0
bath           0
price          0
Price_per_sqft 1265
dtype: int64
```

```
In [40]: df.dropna(inplace = True)
```

```
In [41]: df.isnull().sum()
```

```
Out[41]: location          0
size          0
total_sqft      0
bath           0
price          0
Price_per_sqft  0
dtype: int64
```

```
In [42]: df
```

Out[42]:

	location	size	total_sqft	bath	price	Price_per_sqft
0	Electronic City Phase II	2	1056.0	2.0	39.07	0.036998
1	Chikka Tirupathi	4	2600.0	5.0	120.00	0.046154
2	Uttarahalli	3	1440.0	2.0	62.00	0.043056
3	Lingadheeranahalli	3	1521.0	3.0	95.00	0.062459
4	Kothanur	2	1200.0	2.0	51.00	0.042500
...	...	...	...	...	...	...
13315	Whitefield	5	3453.0	4.0	231.00	0.066898
13316	Richards Town	4	3600.0	5.0	400.00	0.111111
13317	Raja Rajeshwari Nagar	2	1141.0	2.0	60.00	0.052585
13318	Padmanabhanagar	4	4689.0	4.0	488.00	0.104073
13319	Doddathoguru	1	550.0	1.0	17.00	0.030909

11935 rows × 6 columns

In [43]:

```
x = df[['size','total_sqft', 'bath','Price_per_sqft' ]]  
y = df['price']
```

In [44]:

```
x
```

Out[44]:

	size	total_sqft	bath	Price_per_sqft
0	2	1056.0	2.0	0.036998
1	4	2600.0	5.0	0.046154
2	3	1440.0	2.0	0.043056
3	3	1521.0	3.0	0.062459
4	2	1200.0	2.0	0.042500
...	...	...	...	...
13315	5	3453.0	4.0	0.066898
13316	4	3600.0	5.0	0.111111
13317	2	1141.0	2.0	0.052585
13318	4	4689.0	4.0	0.104073
13319	1	550.0	1.0	0.030909

11935 rows × 4 columns

In [45]:

```
x_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=42 )
```

In [46]:

```
model = LinearRegression()
```

```
In [47]: model.fit(X_train, y_train)
```

```
Out[47]: ▼ LinearRegression  
LinearRegression()
```

```
In [48]: y_pred = model.predict(X_test)
```

```
In [49]: mse = mean_squared_error(y_test, y_pred)
```

```
In [50]: r_squared = r2_score(y_test, y_pred)
```

```
In [51]: cv = np.mean(cross_val_score(model, X, y, cv=5))
```

```
In [52]: print(f'Mean Squared Error (MSE): {mse}')
```

```
print(f'R-squared: {r_squared}')
```

```
print(f'Cross Validation Score: {cv}')
```

Mean Squared Error (MSE): 1131.4300760113033

R-squared: 0.7793564127656505

Cross Validation Score: 0.6983224679434117

```
In [ ]:
```